

# CNN to Classify American Sign Language (ASL)

...

Group 9

Toufik Bouras, Noah Olsen, Teryn Zmuda



# Overview

Problem: Classify American sign language images to letters of the alphabet

Method: Apply CNN with Pytorch, and vary approach with three network types, including VGG-16, Resnet.

Data: Image data set for the American alphabet sign language, 1 GB


- Training: 87,000 images, 64 x 64 pixels
  - Classify into 29 classes, 26 letters A-Z and 3 for space, delete, and nothing
  - 3,000 files for each class
- Validation: 29 images, 64x64 pixels

# **Convolutional Neural Network (CNN): Multi-layer perceptron**

# Method Specifics & Theory

- Mini-batch gradient descent
- 29 classes
- Classification organization

```
input_size = 784
hidden_size = 50000
num_classes = 29
num_epochs = 5
batch_size = 100
learning_rate = 0.001
```



```
classes = ('A', 'B', 'C', 'D', 'del', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N',  
          'nothing', 'O', 'P', 'Q', 'R', 'S', 'space', 'T', 'U', 'V', 'W', 'X', 'Y', 'Z')
```

# Data Challenges

- Extremely large dataset.
- The names of the classes are based on the folder names of each image.
- Use the Pytorch function ImageFolder to handle this kind of data.

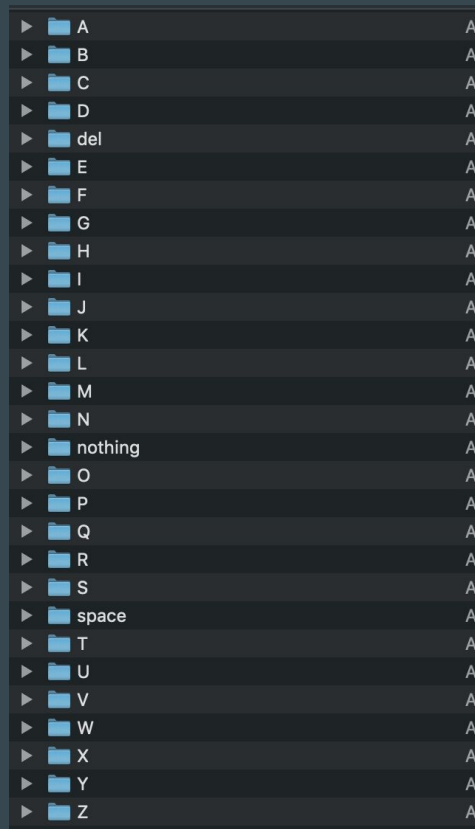
```
# Data augmentation and normalization for training
# Just normalization for validation
data_transforms = {
    'asl_alphabet_train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    'asl_alphabet_test': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}

image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
                                          data_transforms[x])
                  for x in ['asl_alphabet_train', 'asl_alphabet_test']}


dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=batch_size,
                                              shuffle=True)
               for x in ['asl_alphabet_train', 'asl_alphabet_test']}

dataset_sizes = {x: len(image_datasets[x]) for x in ['asl_alphabet_train', 'asl_alphabet_test']}

class_names = image_datasets['asl_alphabet_train'].classes
```



# Training Images for “A”



Data augmentation and  
normalization for training


Normalization for validation

```
data_transforms = {  
    'asl_alphabet_train': transforms.Compose([  
        transforms.Grayscale(1),  
        transforms.RandomResizedCrop(224),  
        transforms.RandomHorizontalFlip(),  
        transforms.ToTensor(),  
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])  
    ]),  
    'asl_alphabet_test': transforms.Compose([  
        transforms.Grayscale(1),  
        transforms.Resize(256),  
        transforms.CenterCrop(224),  
        transforms.ToTensor(),  
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])  
    ]),  
}
```

# Understanding the different networks



# VGG-16

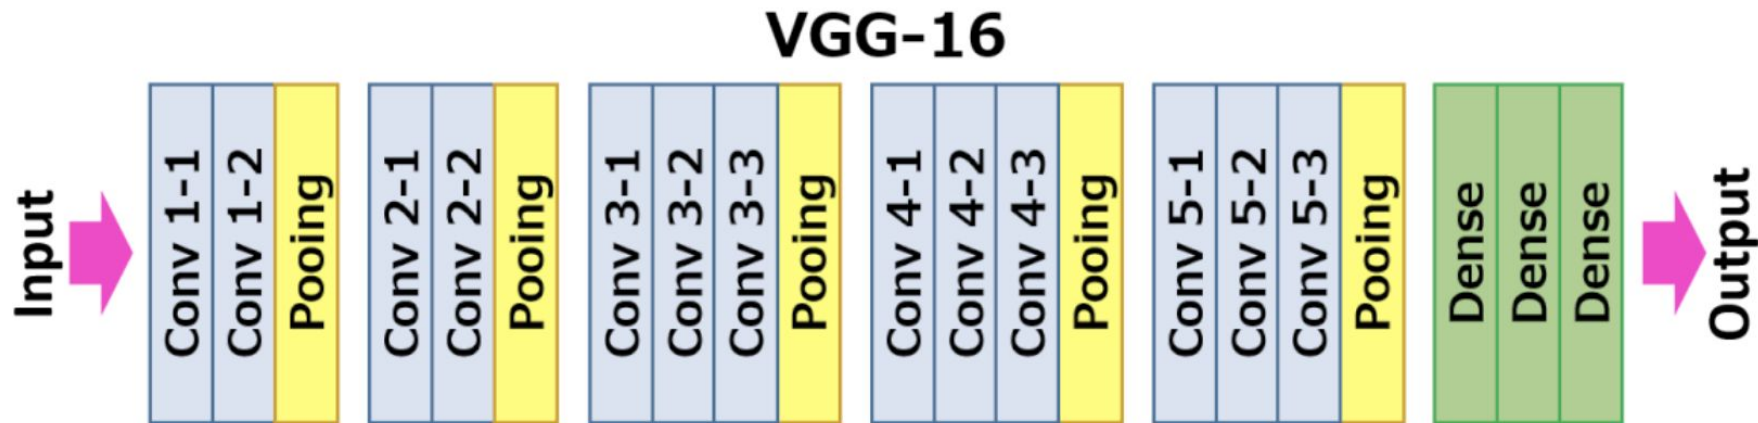


VGG16

## VGG-16:

- 2014 genesis, build from AlexNet, Oxford
- Competitor: GoogleNet
- Deeper but smaller filters, only 3 x 3 CONV all the way: 8 layers (AlexNet) to 16 layers
- Total memory = 100 MB/Image (forward only); 138M parameters

# VGG-16 Network Diagram




# VGG-16 Network Code

```
# Choose the right argument for x
net = models.vgg16(pretrained=True)
# Freeze model weights
for param in net.parameters():
    param.requires_grad = False

num_fters = net.classifier[6].in_features

# Add on classifier
net.classifier[6] = nn.Sequential(
    nn.Linear(num_fters, 256),
    nn.ReLU(),
    nn.Dropout(0.4),
    nn.Linear(256, num_classes),
    nn.LogSoftmax(dim=1))
net.cuda()
net = nn.DataParallel(net)
#
```


# Resnet



Resnet:

- Follow up on Alexnet
- Stack resenet blocks
- Batch normalization after each layer
- Deeper than previous networks
- Residual connections
- Can train thousands of layers and still achieve good performance.

# Resnet Network Diagram



# Resnet Network Code

```
#-----
net = models.resnet50(pretrained=True)
# Freeze model weights
for param in net.parameters():
    param.requires_grad = False

num_fts = net.fc.in_features

#Add on classifier
net.fc = nn.Sequential(
    nn.Linear(num_fts, 256),
    nn.ReLU(),
    nn.Dropout(0.4),
    nn.Linear(256, num_classes),
    nn.LogSoftmax(dim=1))

net.cuda()
net = nn.DataParallel(net)
#-----
criterion = nn.NLLLoss();
optimizer = torch.optim.Adam(net.parameters(), lr=learning_rate);

j = 0
loss_list = []
iteration_list = []
accuracy_list = []
```

# Our Model

Our attempt at a more custom fitted model. The initial inspiration was Resnet.

Features:

- Deeper than previous networks
- Residual connections

---

# Code

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()

        self.conv1 = torch.nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
        self.conv1_bn = nn.BatchNorm2d(32)
        self.relu = nn.ReLU()
        self.conv2 = torch.nn.Conv2d(32, 32, kernel_size=3, stride=1, padding=1)
        self.conv2_bn = nn.BatchNorm2d(32)
        self.relu = nn.ReLU()
        self.pool = torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        self.conv3 = torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.conv3_bn = nn.BatchNorm2d(64)
        self.relu = nn.ReLU()
        self.conv4 = torch.nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1)
        self.conv4_bn = nn.BatchNorm2d(64)
        self.relu = nn.ReLU()
        self.pool = torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=0)

        self.conv5 = torch.nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1)
        self.conv5_bn = nn.BatchNorm2d(128)
        self.conv6 = torch.nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1)
        self.conv6_bn = nn.BatchNorm2d(128)
        self.relu = nn.ReLU()

        self.fc1 = torch.nn.Linear(128 * 28 * 28, 128, bias=True)
        self.relu = nn.ReLU()
        self.dropout = nn.Dropout(0.2)
        self.out = nn.Linear(128, 29, bias=True)
        self.softmax = nn.LogSoftmax(dim=1)
```

```
def forward(self, x):
    out = self.conv1(x)
    out = self.conv1_bn(out)
    out = F.relu(out)
    out = F.relu(self.conv2_bn(self.conv2(out)))
    out = self.pool(out)

    out = F.relu(self.conv3_bn(self.conv3(out)))
    out = F.relu(self.conv4_bn(self.conv4(out)))
    out = self.pool(out)

    out = self.relu(self.conv5_bn(self.conv5(out)))
    out = self.relu(self.conv6_bn(self.conv6(out)))
    out = self.pool(out)

    #print(out.shape)

    out = out.view(-1, 128 * 28 * 28)

    out = self.dropout(F.relu(self.fc1(out)))
    out = self.softmax(self.out(out))

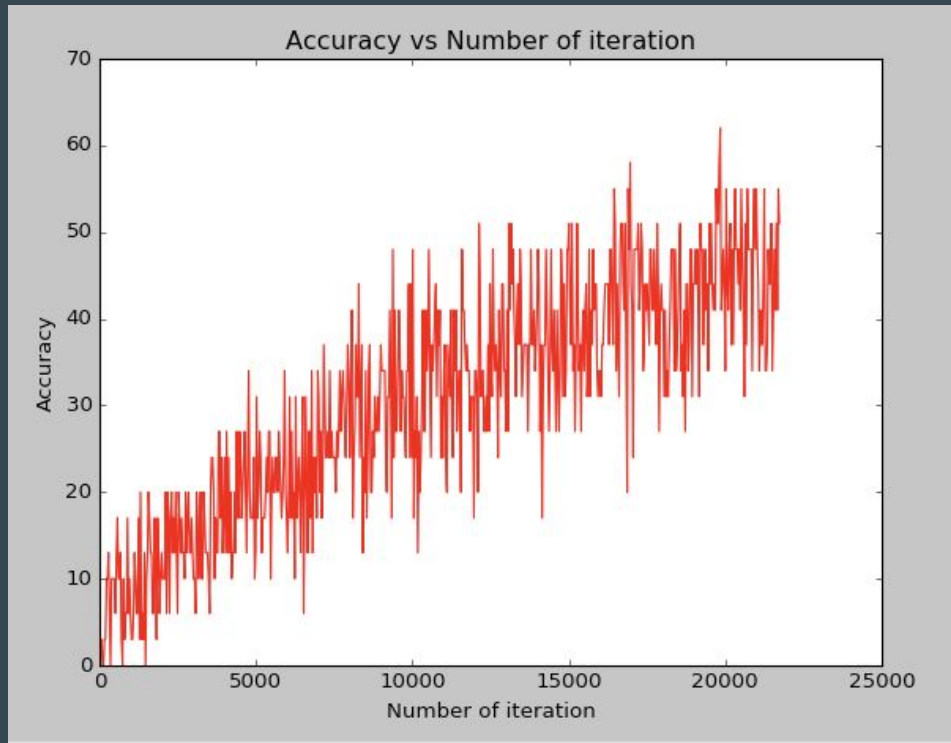
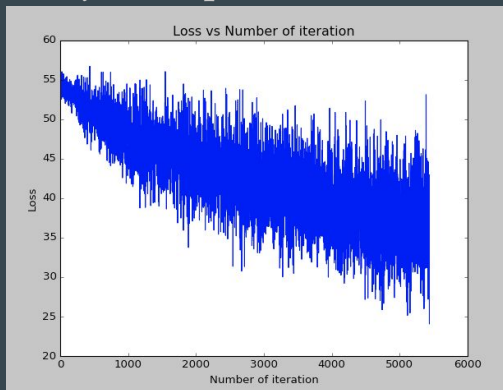
    return (out)
```



# Accuracy and Loss Over Time for our Model

## Results:



- Started off with very low accuracy
- Learned significantly over the lifetime of the model.
- Loss improved over time.
- Very slow performance



# Accuracy and Loss Over Time for adapted VGG16

## Results:



- Very slow.
- Performance is quite strong.
- Finished with around 80% accuracy.



# Accuracy and Loss Over Time for adapted Resnet50

## Results:

- Started off with very low accuracy
- Scaled rapidly
- Impressively quick performance
- Finished with around 90% accuracy



Questions?

# Sources

- <https://neurohive.io/en/popular-networks/vgg16/>
- <https://arxiv.org/abs/1512.03385>
- Akash. (2018, April 22). ASL Alphabet. Retrieved from <https://www.kaggle.com/grassknotted/asl-alphabet>
- Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. CoRR, abs/1409.1556 . Retrieved from <http://arxiv.org/abs/1409.1556>
- Stanford School of Engineering, CNN Architecture Lecture. Aug. 11, 2017. Retrieved from <https://www.youtube.com/watch?v=DAOcjjcFr1>